

# SOCI 40258

Causal Mediation Analysis

Week 9: Robust Estimation

# Outline

- Model misspecification
- Robust/DML estimation of total effects
- Robust/DML estimation of natural effects
- Unobserved confounding
- Sensitivity analysis

# The problem of misspecification

- All the parametric approaches to estimation that we've covered thus far hinge on the assumption of no model misspecification
- In practice, however, model misspecification is likely, if not certain, which will lead to bias and inconsistency
- Triangulating results from different models and estimation procedures can help to allay concerns about misspecification, but standard inferential statistics do not properly account for this specification search

# Multiply robust estimation

- Most of the estimators we have covered require correct parametric models for the outcome, the mediator(s), and/or the exposure
- Multiply robust estimators only require correct models for a subset of these variables in order to remain consistent
  - When implemented with parametric models, multiply robust estimators provide a degree of protection against misspecification bias
  - When implemented with data-adaptive machine learning (ML) methods, multiply robust estimators provide even greater protection against problems associated with model misspecification

# Robust estimation of total effects

- We will begin by reviewing on robust estimation of average total effects in order to establish foundational principles
- In particular, we will focus on constructing a robust estimator for the ATE using its efficient influence function
- The resulting estimator will be doubly robust, that is, consistent if either a model for the exposure or a model for the outcome is correctly specified

# Efficient Influence Functions (EIFs)

- An efficient influence function provides a mechanism for understanding the best possible performance that an estimator can achieve under certain conditions
- In general, an influence function captures how a single data point affects the overall estimate of a given target parameter (e.g., the ATE)
- An efficient influence function, then, is the influence function corresponding to an estimator that converges the lowest possible asymptotic variance
  - Essentially, it represents the optimal way that an estimator can be influenced by the data

# EIF for the ATE

- With a binary treatment  $D \in \{0,1\}$  and a set of baseline confounders  $C$ , the EIF for the average total effect,  $ATE(1,0) = E(Y(1) - Y(0))$ , can be expressed as follows:

$$EIF_{ATE} = \frac{1}{n} \sum \left[ \left[ \frac{D}{\pi_1(C)} - \frac{1-D}{1-\pi_1(C)} \right] \times (Y - \mu_D(C)) \right] + [\mu_1(C) - \mu_0(C)] - ATE(1,0)$$

where  $\pi_d(C) = P(D = d|C)$ ,  $\mu_d(C) = E(Y|C, D = d)$ , and  $Y - \mu_D(C) = Y - E(Y|C, D)$  is a residual term for the outcome

# Robust estimation for the ATE

- Setting the  $EIF_{ATE}$  equal to zero and solving yields an efficient, doubly-robust (DR) estimator for the ATE:

$$\widehat{ATE}(1,0)^{dr} = \frac{1}{n} \sum \left[ \left[ \frac{D}{\hat{\pi}_1(C)} - \frac{1-D}{1-\hat{\pi}_1(C)} \right] \times (Y - \hat{\mu}_D(C)) \right] + [\hat{\mu}_1(C) - \hat{\mu}_0(C)]$$

where the hats are used to denote estimates of the conditional means and probabilities



# Robust estimation for the ATE

- Setting the  $EIF_{ATE}$  equal to zero and solving yields an efficient, doubly-robust (DR) estimator for the ATE:

$$\widehat{ATE}(1,0)^{dr} = \frac{1}{n} \sum \left[ \underbrace{\left[ \frac{D}{\hat{\pi}_1(C)} - \frac{1-D}{1-\hat{\pi}_1(C)} \right] \times (Y - \hat{\mu}_D(C))}_{\text{IPW update based on residual confounding}} + \underbrace{[\hat{\mu}_1(C) - \hat{\mu}_0(C)]}_{\text{regression imputation}} \right]$$

# Robust estimation for the ATE with parametric models

- With a DR estimator for the ATE,

$$\widehat{ATE}(1,0)^{dr} = \frac{1}{n} \sum \left[ \left[ \frac{D}{\hat{\pi}_1(C)} - \frac{1-D}{1-\hat{\pi}_1(C)} \right] \times (Y - \hat{\mu}_D(C)) \right] + [\hat{\mu}_1(C) - \hat{\mu}_0(C)],$$

estimation with parametric models just involves fitting GLMs for  $P(D|C)$  and  $E(Y|C, D)$  and then using these models to compute all the terms in  $\widehat{ATE}(1,0)^{dr}$

- If either the GLM for  $P(D|C)$  or the GLM for  $E(Y|C, D)$  is correctly specified, this estimator will be consistent and asymptotically normal, providing a degree of protection against model misspecification

# Robust estimation for the ATE with parametric models

- The DR estimator can be implemented with parametric models through the following series of steps:
  1. Fit a GLM for  $P(D|C)$ 
    - Use the fitted model to compute  $\hat{\pi}_1(C)$
  2. Fit a GLM for  $E(Y|C, D)$ 
    - Use the fitted model to compute  $\hat{\mu}_1(C)$ ,  $\hat{\mu}_0(C)$ , and  $Y - \hat{\mu}_D(C)$
  3. Plug these terms into  $\widehat{ATE}(1,0)^{dr}$  for each sample member and solve

# Robust estimation for the ATE with ML models

- With a DR estimator for the ATE,

$$\widehat{ATE}(1,0)^{dr} = \frac{1}{n} \sum \left[ \left[ \frac{D}{\hat{\pi}_1(C)} - \frac{1-D}{1-\hat{\pi}_1(C)} \right] \times (Y - \hat{\mu}_D(C)) \right] + [\hat{\mu}_1(C) - \hat{\mu}_0(C)],$$

estimation via ML just involves training ML models for  $P(D|C)$  and  $E(Y|C, D)$ , and then using these models to compute all the terms in  $\widehat{ATE}(1,0)^{dr}$

- If all the ML models possess sufficiently fast rates of convergence, then the DR estimator will be consistent, efficient, and asymptotically normal, providing additional protection against misspecification
  - Generally, each of the ML models must be at least  $n^{1/4}$ -consistent (i.e., they must converge to their target estimand at a  $n^{1/4}$  rate)

# Robust estimation for the ATE with ML models

- The DR estimator can be implemented with ML models through the following series of steps:
  1. Train a ML model for  $P(D|C)$ 
    - Use the trained model to compute  $\hat{\pi}_1(C)$
  2. Train a ML model for  $h(Y|C, D)$ 
    - Use the trained model to compute  $\hat{\mu}_1(C)$ ,  $\hat{\mu}_0(C)$ , and  $Y - \hat{\mu}_D(C)$
  3. Plug these terms into  $\widehat{ATE}(1,0)^{dr}$  for each sample member and solve

# Robust estimation for the ATE with ML models

- Candidate ML models for  $P(D|C)$  and  $E(Y|C, D)$  might include:
  - Regularized GLMs (LASSO, ridge, elasticnet)
  - Classification and regression trees (CARTs)
  - Random forests (RFs)
  - Gradient boosted trees (GBTs)
  - Artificial neural networks (ANNs)
  - Super learners (SLs)

# Sample splitting

- With ML models, using the same data for training the models and for estimating the ATE can sometimes lead to bias and complicates standard approaches to statistical inference
- To avoid these problems, we can pair DR estimation with a sample splitting algorithm
- The sample splitting algorithm proceeds as follows:
  1. Split the original sample into two separate subsamples
  2. Train the ML models for  $P(D|C)$  and  $E(Y|C, D)$  using the first subsample
  3. Apply the trained models to the second subsample to compute the  $\widehat{ATE}(1,0)^{\text{dr}}$

# Repeated cross-fitting

- Sample splitting is inefficient, since each subset of the data is only used for training or effect estimation, but not both
- Cross-fitting resolves this problem by iterating the sample splitting procedure:
  1. Split the original sample into  $4 \leq J \leq 10$  equally sized subsamples
  2. Train the ML models for  $P(D|C)$  and  $E(Y|C, D)$  using  $J - 1$  of the subsamples
  3. Apply the trained models to the remaining subsample  $j$  to compute  $\widehat{ATE}(1,0)_j^{\text{dr}}$
  4. Iterate the previous two steps until each subsample is used for estimation a single time
  5. Compute  $\widehat{ATE}(1,0)^{\text{dr}} = \frac{1}{J} \sum_{j=1}^J \widehat{ATE}(1,0)_j^{\text{dr}}$  by averaging estimates across the subsamples



# Wald tests and CIs

- When computed using cross-fitting, the  $\widehat{ATE}(1,0)^{dr}$  is asymptotically normally distributed with a variance equal to  $Var(EIF_{ATE})/n$
- This allows the use of standard inferential statistics, including...
  - Wald tests, based on  $W = \frac{\widehat{ATE}(1,0)^{dr} - ATE(1,0)_{H_0}}{\sqrt{\widehat{Var}(EIF_{ATE})/n}} \sim N(0,1)$  under  $H_0$
  - Wald confidence intervals, given by  $\widehat{ATE}(1,0)^{dr} \pm Z_\alpha \times \sqrt{\widehat{Var}(EIF_{ATE})/n}$  where  $Z \sim N(0,1)$
- The nonparametric bootstrap can also be used to test null hypotheses and construct confidence intervals for cross-fit ML estimators, but it is more computationally intensive

# Example: NLSY79

- Does attending college reduce depression later in adulthood?
- 1979 National Longitudinal Study of Youth
  - Exposure ( $D$ )
    - sample member attended college before age 22
  - Outcome ( $Y$ ):
    - standardized scores on the CES-D at age 40
  - Covariates ( $C$ ):
    - Race and gender
    - Parental education, occupation, and income
    - Household size
    - AFQT scores in high school

# Example: NLSY79

- Compute doubly robust estimates for the *ATE* using parametric models

```
1  ### wk 9 nlsy tutorial ###
2  rm(list=ls())
3
4  # load/install libraries #
5  packages<-c("dplyr", "tidyr", "foreign", "margins", "survey", "ranger", "xgboost", "SuperLearner")
6  #install.packages(packages)
7
8  for (package.i in packages) {
9    suppressPackageStartupMessages(library(package.i, character.only=TRUE))
10  }
11
12  # load data #
13  datadir <- "C:/Users/Geoffrey Wodtke/Dropbox/D/courses/2023-24_UOFCHICAGO/SOCI_40258_CAUSAL_MEDIAT
14  nlsy <- read.dta(paste(datadir, "nlsy79.dta", sep=""))
15
16  nlsy <- nlsy[complete.cases(nlsy[, c("cesd_age40", "ever_unemp_age3539", "att22",
17    "female", "black", "hispan", "paredu", "parprof", "parinc_prank", "famsize", "afqt3")]),]
18
19  nlsy$std_cesd_age40 <- (nlsy$cesd_age40-mean(nlsy$cesd_age40))/sd(nlsy$cesd_age40)
20
21  # parametric doubly robust estimate for ATE #
22  Ymodel.lm <- lm(std_cesd_age40 ~ att22*(female + black + hispan + paredu + parprof +
23    parinc_prank*parinc_prank + famsize + afqt3), data=nlsy)
24
25  Dmodel.pr <- glm(att22 ~ female + black + hispan + paredu + parprof +
26    parinc_prank + famsize + afqt3, data=nlsy, family=binomial(link="probit"))
27
```

# Example: NLSY79

- Compute doubly robust estimates for the *ATE* using parametric models

```
31 gdata <- nlsy
32
33 gdata$att22 <- 1
34 yhat1 <- predict(Ymodel.lm, newdata=gdata)
35
36 gdata$att22 <- 0
37 yhat0 <- predict(Ymodel.lm, newdata=gdata)
38
39 resid <- Ymodel.lm$residuals
40
41 phatD1 <- predict(Dmodel.pr, type = "response")
42
43 ipw <- (nlsy$att22*mean(nlsy$att22)/phatD1) - (1-nlsy$att22)*(1-mean(nlsy$att22))/(1-phatD1)
44
45 eif <- ipw*resid + (yhat1-yhat0)
```

# Example: NLSY79

- Compute doubly robust estimates for the *ATE* using parametric models

```
46 ATEhat <- mean(eif)
47
48 VarATEhat <- var(eif)/length(ipw)
49
50 CI95pct <- c(ATEhat-1.96*sqrt(VarATEhat), ATEhat+1.96*sqrt(VarATEhat))
51
52 p <- (1-pnorm(abs(ATEhat/sqrt(VarATEhat)),0,1))*2
53
54 parResults <- data.frame(
55   point.est=round(ATEhat, digits=3),
56   se=round(sqrt(VarATEhat), digits=3),
57   ll.95ci=round(CI95pct[1], digits=3),
58   ul.95ci=round(CI95pct[2], digits=3),
59   pval=round(p, digits=3))
60 rownames(parResults) <- c('ATE')
61
62 print(parResults)
```

```
> print(parResults)
      point.est      se ll.95ci ul.95ci pval
ATE      -0.15 0.022  -0.194  -0.107    0
> |
```

# Example: NLSY79

- Compute DML estimates for the *ATE* using super learners

```
64 # semi-parametric doubly robust estimate for ATE w/ cross-fitting #
65 set.seed(3308004)
66
67 nlsy$k <- runif(nrow(nlsy), 0, 1)
68 nlsy <- nlsy[order(nlsy$k), ]
69 nlsy$k <- rep(1:5, length.out=nrow(nlsy))
70
71 cntrl.sl <- SuperLearner.CV.control(V=10)
72
73 confounders <- c("female", "black", "hispan", "paredu",
74                 "parprof", "parinc_prank", "famsize", "afqt3")
```

# Example: NLSY79

- Compute DML estimates for the *ATE* using super learners

```
76 for (j in 1:5) {  
77  
78   df.train <- nlsy[which(nlsy$k!=j),]  
79   df.est <- nlsy[which(nlsy$k==j),]  
80  
81   Ymodel.sl <- SuperLearner(  
82     Y=df.train[, "std_cesd_age40"],  
83     X=df.train[, c("confounders", "att22")],  
84     SL.library=c("SL.lm", "SL.ranger", "SL.xgboost"),  
85     cvControl=cntrl.sl)  
86  
87   Dmodel.sl <- SuperLearner(  
88     Y=df.train[, "att22"],  
89     X=df.train[, "confounders"],  
90     family = binomial(),  
91     SL.library=c("SL.glm", "SL.ranger", "SL.xgboost"),  
92     cvControl=cntrl.sl)
```

# Example: NLSY79

- Compute DML estimates for the *ATE* using super learners

```
94 gdata <- df.est[,c(confounders, "att22")]
95
96 gdata$att22 <- 1
97 nlsy$yhat1[nlsy$k==j] <- predict(Ymodel.sl, newdata=gdata)$pred
98
99 gdata$att22 <- 0
100 nlsy$yhat0[nlsy$k==j] <- predict(Ymodel.sl, newdata=gdata)$pred
101
102 nlsy$resid[nlsy$k==j] <- df.est$std_cesd_age40 -
103   predict(Ymodel.sl, df.est[,c(confounders, "att22")])$pred
104
105 phatD1 <- predict(Dmodel.sl, newdata=df.est[,confounders], type="prob")$pred
106
107 nlsy$ipw[nlsy$k==j] <- (nlsy$att22[nlsy$k==j]/phatD1)
108   - (1-nlsy$att22[nlsy$k==j])/(1-phatD1)
109 }
```



# Example: NLSY79

- Compute DML estimates for the *ATE* using super learners

```
111 eif <- nlsy$ipw*nlsy$resid + (nlsy$yhat1-nlsy$yhat0)
112
113 ATEhat <- mean(eif)
114
115 VarATEhat <- var(eif)/length(nlsy$ipw)
116
117 CI95pct <- c(ATEhat-1.96*sqrt(VarATEhat), ATEhat+1.96*sqrt(VarATEhat))
118
119 p <- (1-pnorm(abs(ATEhat/sqrt(VarATEhat)),0,1))*2
120
121 sparResults <- data.frame(
122   point.est=round(ATEhat, digits=3),
123   se=round(sqrt(VarATEhat), digits=3),
124   ll.95ci=round(CI95pct[1], digits=3),
125   ul.95ci=round(CI95pct[2], digits=3),
126   pval=round(p, digits=3))
127 rownames(sparResults) <- c('ATE')
128
129 print(sparResults)
```

```
> print(sparResults)
      point.est      se ll.95ci ul.95ci  pval
ATE      -0.117 0.046  -0.208  -0.026 0.012
```

# Robust estimation for natural effects

- Next, we will focus on multiply robust estimation of natural direct and indirect effects based on their efficient influence functions...

# Robust estimation for natural effects

- Define the marginal means of the nested potential outcomes as follows:

$$\psi(d_1, d_2) = E\left(Y(d_1, M(d_2))\right),$$

where  $d_1$  and  $d_2$  just denote two generic values of the exposure

- With this notation, then, we can define our target estimands as follows:

$$NDE(d, d^*) = \psi(d, d^*) - \psi(d^*, d^*) = E\left(Y(d, M(d^*))\right) - E\left(Y(d^*, M(d^*))\right)$$

$$NIE(d, d^*) = \psi(d, d) - \psi(d, d^*) = E\left(Y(d, M(d))\right) - E\left(Y(d, M(d^*))\right)$$

# The EIF for $\psi_{d_1, d_2}$

- With a binary treatment  $D \in \{0, 1\}$  and set of baseline confounders  $C$ , the EIF for the marginal mean of the nested potential outcomes can be expressed as:

$$\begin{aligned} EIF_{\psi(d_1, d_2)} = & \frac{1}{n} \sum \left[ \left( \frac{I(D=d_1)}{\pi_{d_2}(C)} \right) \left( \frac{\pi_{d_2}(C, M)}{\pi_{d_1}(C, M)} \right) (Y - \mu_{d_1}(C, M)) \right] + \\ & \left[ \left( \frac{I(D=d_1)}{\pi_{d_2}(C)} \right) \left( \frac{\pi_{d_2}(C, M)}{\pi_{d_1}(C, M)} \right) (\mu_{d_1}(C, M) - \nu_{d_2}(C)) \right] + \\ & \nu_{d_2}(C) - \psi(d_1, d_2) \end{aligned}$$

where  $\pi_{d_2}(C) = P(D = d_2|C)$ ,  $\pi_{d_2}(C, M) = P(D = d_2|C, M)$  and  $\pi_{d_1}(C, M)$  is defined analogously,  $\mu_{d_1}(C, M) = E(Y|C, D = d_1, M)$ , and  $\nu_{d_2}(C) = E(E(Y|C, D = d_1, M)|C, D = d_2)$

# Robust estimation for natural effects

- Setting the  $EIF_{\psi_{d_1, d_2}}$  equal to zero and solving yields an efficient, multiply robust (MR) estimator for the mean of the nested potential outcomes:

$$\begin{aligned}\hat{\psi}(d_1, d_2)^{\text{mr}} = & \frac{1}{n} \sum \left[ \left( \frac{I(D=d_1)}{\hat{\pi}_{d_2}(C)} \right) \left( \frac{\hat{\pi}_{d_2}(C, M)}{\hat{\pi}_{d_1}(C, M)} \right) (Y - \hat{\mu}_{d_1}(C, M)) \right] + \\ & \left[ \left( \frac{I(D=d_1)}{\hat{\pi}_{d_2}(C)} \right) \left( \frac{\hat{\pi}_{d_2}(C, M)}{\hat{\pi}_{d_1}(C, M)} \right) (\hat{\mu}_{d_1}(C, M) - \hat{v}_{d_2}(C)) \right] + \\ & \hat{v}_{d_2}(C)\end{aligned}$$

where the hats are used to denote estimates of the conditional means and probabilities

# Robust estimation for natural effects

- Comparing different estimated marginal means for the nested potential outcomes yields MR estimators for the natural direct and indirect effects of interest:

$$\widehat{NDE}(1,0)^{\text{mr}} = \hat{\psi}(1,0)^{\text{mr}} - \hat{\psi}(0,0)^{\text{mr}}$$

$$\widehat{NDE}(1,0)^{\text{mr}} = \hat{\psi}(1,1)^{\text{mr}} - \hat{\psi}(1,0)^{\text{mr}}$$

# Robust estimation for natural effects with parametric models

- With the MR estimator for  $\psi(d_1, d_2)$ , estimation with parametric models involves fitting GLMs for  $P(D|C)$ ,  $P(D|C, M)$ ,  $E(Y|C, D, M)$ , and  $E(E(Y|C, D = d_1, M)|C, D)$ 
  - Then, these models are used to compute all the terms in  $\hat{\psi}(d_1, d_2)^{\text{mr}}$
- If the models for (i)  $P(D|C)$  and  $P(D|C, M)$ , (ii) for  $P(D|C)$  and  $E(Y|C, D, M)$ , or (iii) for  $E(Y|C, D, M)$  and  $E(E(Y|C, D = d_1, M)|C, D)$  are correctly specified, this estimator will be consistent, efficient, and asymptotically normal
  - It is sometimes said to be “triply robust” because there are three distinct pathways to consistency

# Robust estimation for natural effects with parametric models

- The MR estimator can be implemented with parametric models through the following series of steps:
  1. Fit a GLM for  $P(D|C)$ 
    - Use the fitted model to compute  $\hat{\pi}_{d_2}(C)$
  2. Fit a GLM for  $P(D|C, M)$ 
    - Use the fitted model to compute  $\hat{\pi}_{d_1}(C, M)$  and  $\hat{\pi}_{d_2}(C, M)$
  3. Fit a GLM for  $E(Y|C, D, M)$ 
    - Use the fitted model to compute  $\hat{\mu}_{d_1}(C, M)$
  4. Fit a GLM for  $E(E(Y|C, D = d_1, M)|C, D)$ 
    - Use the fitted model to compute  $\hat{v}_{d_2}(C)$
  5. Plug these terms into  $\hat{\psi}(1,1)^{\text{mr}}$ ,  $\hat{\psi}(0,0)^{\text{mr}}$ , and  $\hat{\psi}(1,0)^{\text{mr}}$  and solve



# Robust estimation for natural effects with ML models

- With the MR estimator for  $\psi(d_1, d_2)$ , estimation via ML just involves training ML models for  $P(D|C)$ ,  $P(D|C, M)$ ,  $E(Y|C, D, M)$ , and  $E(E(Y|C, D = d_1, M)|C, D)$ , and then using these models to compute  $\hat{\psi}(d_1, d_2)^{\text{mr}}$
- If all the ML models possess sufficiently fast rates of convergence—generally, meaning they are at least  $n^{1/4}$ -consistent—then the MR estimator will be consistent, efficient, and asymptotically normal

# Robust estimation for natural effects with ML models

- The MR estimator can be implemented with ML models through the following series of steps:
  1. Train a ML model for  $P(D|C)$ 
    - Use the trained model to compute  $\hat{\pi}_{d_2}(C)$
  2. Train a ML model for  $P(D|C, M)$ 
    - Use the trained model to compute  $\hat{\pi}_{d_1}(C, M)$  and  $\hat{\pi}_{d_2}(C, M)$
  3. Train a ML model for  $E(Y|C, D, M)$ 
    - Use the trained model to compute  $\hat{\mu}_{d_1}(C, M)$
  4. Train a ML model for  $E(E(Y|C, D = d_1, M)|C, D)$ 
    - Use the trained model to compute  $\hat{v}_{d_2}(C)$
  5. Plug these terms into  $\hat{\psi}(1,1)^{\text{mr}}$ ,  $\hat{\psi}(0,0)^{\text{mr}}$ , and  $\hat{\psi}(1,0)^{\text{mr}}$  and solve

# Repeated cross-fitting

- With ML models, using the same data for model training and for effect estimation can lead to bias and complicates standard approaches to statistical inference
- To avoid these problems, we can pair the estimation approach outlined previously with the cross-fitting algorithm:
  1. Split the original sample into  $4 \leq J \leq 10$  equally sized subsamples
  2. Train the ML models using  $J - 1$  of the subsamples
  3. Apply the trained models to the remaining subsample  $J$  to compute  $\hat{\psi}(d_1, d_2)_j^{\text{mr}}$
  4. Iterate the previous two steps until each subsample is used for estimation a single time
  5. Compute  $\hat{\psi}(d_1, d_2)^{\text{mr}} = \frac{1}{J} \sum_{j=1}^J \hat{\psi}(d_1, d_2)_j^{\text{mr}}$  by averaging across the  $J$  subsamples

# Wald tests and CIs for natural effects

- When computed using cross-fitting, the  $\widehat{NDE}(1,0)^{\text{mr}}$  is asymptotically normally distributed with a variance equal to  $Var(EIF_{\psi(1,0)} - EIF_{\psi(0,0)})/n$
- Similarly, the  $\widehat{NIE}(1,0)^{\text{mr}}$  is also asymptotically normal with a variance equal to  $Var(EIF_{\psi(1,1)} - EIF_{\psi(1,0)})/n$
- This suggests the use of standard inferential statistics, including Wald tests and confidence intervals, with the standard errors for  $\widehat{NDE}(1,0)^{\text{mr}}$  and  $\widehat{NIE}(1,0)^{\text{mr}}$  given by:

$$se(\widehat{NDE}(1,0)^{\text{mr}}) = \sqrt{\widehat{Var}(EIF_{\psi(1,0)} - EIF_{\psi(0,0)})/n}$$

$$se(\widehat{NIE}(1,0)^{\text{mr}}) = \sqrt{\widehat{Var}(EIF_{\psi(1,1)} - EIF_{\psi(1,0)})/n}$$

# Example: NLSY79

- 1979 National Longitudinal Study of Youth
  - Exposure ( $D$ )
    - sample member attended college before age 22
  - Outcome ( $Y$ )
    - standardized scores on the CES-D at age 40
  - Covariates ( $C$ )
    - Race, gender, parental education, occupation, and income, household size, AFQT scores
  - A potential mediator ( $M$ )
    - unemployment between age 35-40

# Example: NLSY79

- Many studies have documented that going to college seems to reduce the likelihood of becoming depressed later in life—but how does this effect come about?
- One possibility is that a more advanced education reduces depression by protecting its recipients from financially strenuous and mentally taxing spells of unemployment
- Does unemployment mediate the effect of college attendance on depression?

# Example: NLSY79

- Compute triply robust estimates of the *NDE* and *NIE* using parametric models

```
132 # parametric multiply robust estimate for NDE and NIE #
133 Ymodel.lm <- lm(std_cesd_age40 ~ (att22 * ever_unemp_age3539) * (female + black + hispan +
134   paredu + parprof + parinc_prank*parinc_prank + famsize + afqt3), data=nlsy)
135
136 Dmodel_1.pr <- glm(att22 ~ female + black + hispan + paredu + parprof +
137   parinc_prank + famsize + afqt3, data=nlsy, family=binomial(link="probit"))
138
139 Dmodel_2.pr <- glm(att22 ~ ever_unemp_age3539 + female + black + hispan + paredu + parprof +
140   parinc_prank + famsize + afqt3, data=nlsy, family=binomial(link="probit"))
141
142 phat_d1_C <- predict(Dmodel_1.pr, type = "response")
143 phat_d0_C <- 1-phat_d1_C
144
145 phat_d1_CM <- predict(Dmodel_2.pr, type = "response")
146 phat_d0_CM <- 1-phat_d1_CM
147
148 phat_d1 <- mean(nlsy$att22)
149 phat_d0 <- 1-phat_d1
150
151 f1 <- nlsy$att22*phat_d1 / phat_d1_C
152 f2 <- (1-nlsy$att22)*phat_d0 / phat_d0_C
153 f3 <- nlsy$att22*phat_d0 / phat_d0_C
154
155 s1 <- phat_d0_CM / phat_d1_CM
156
157 resid <- Ymodel.lm$residuals
```

# Example: NLSY79

- Compute triply robust estimates of the *NDE* and *NIE* using parametric models

```
159 gdata <- nlsy
160
161 gdata$att22 <- 1
162 qhat1 <- predict(Ymodel.lm, newdata=gdata)
163
164 gdata$att22 <- 0
165 qhat0 <- predict(Ymodel.lm, newdata=gdata)
166
167 qhat0_CD <- lm(qhat0 ~ att22 * (female + black + hispan +
168   paredu + parprof + parinc_prank + famsize + afqt3), data=nlsy)
169
170 gdata$att22 <- 0
171 EMgivC0qhat0 <- predict(qhat0_CD, newdata=gdata)
172
173 qhat1_CD <- lm(qhat1 ~ att22 * (female + black + hispan +
174   paredu + parprof + parinc_prank + famsize + afqt3), data=nlsy)
175
176 gdata$att22 <- 1
177 EMgivC1qhat1 <- predict(qhat1_CD, newdata=gdata)
178
179 gdata$att22 <- 0
180 EMgivC0qhat1 <- predict(qhat1_CD, newdata=gdata)
181
182 eif_psi_11 <- f1*1*resid + f1*(qhat1 - EMgivC1qhat1) + EMgivC1qhat1
183 eif_psi_00 <- f2*1*resid + f2*(qhat0 - EMgivC0qhat0) + EMgivC0qhat0
184 eif_psi_10 <- f3*s1*resid + f3*s1*(qhat1 - EMgivC0qhat1) + EMgivC0qhat1
```



# Example: NLSY79

- Compute triply robust estimates of the *NDE* and *NIE* using parametric models

```
186 NDEhat <- mean(eif_psi_10 - eif_psi_00)
187 NIEhat <- mean(eif_psi_11 - eif_psi_10)
188 ATEhat <- mean(eif_psi_11 - eif_psi_00)
189
190 VarNDEhat <- var(eif_psi_10 - eif_psi_00)/length(resid)
191 VarNIEhat <- var(eif_psi_11 - eif_psi_10)/length(resid)
192 VarATEhat <- var(eif_psi_11 - eif_psi_00)/length(resid)
193
194 NDE95pctCI <- c(NDEhat-1.96*sqrt(VarNDEhat), NDEhat+1.96*sqrt(VarNDEhat))
195 NIE95pctCI <- c(NIEhat-1.96*sqrt(VarNIEhat), NIEhat+1.96*sqrt(VarNIEhat))
196 ATE95pctCI <- c(ATEhat-1.96*sqrt(VarATEhat), ATEhat+1.96*sqrt(VarATEhat))
197
198 ATEp <- (1-pnorm(abs(ATEhat/sqrt(VarATEhat)),0,1))*2
199 NDEp <- (1-pnorm(abs(NDEhat/sqrt(VarNDEhat)),0,1))*2
200 NIEp <- (1-pnorm(abs(NIEhat/sqrt(VarNIEhat)),0,1))*2
201
202 parMedResults <- data.frame(
203   point.est = round(c(ATEhat, NDEhat, NIEhat), digits = 3),
204   se = round(c(sqrt(VarATEhat), sqrt(VarNDEhat), sqrt(VarNIEhat)), digits = 3),
205   ll.95ci = round(c(ATE95pctCI[1], NDE95pctCI[1], NIE95pctCI[1]), digits = 3),
206   ul.95ci = round(c(ATE95pctCI[2], NDE95pctCI[2], NIE95pctCI[2]), digits = 3),
207   pval = round(c(ATEp, NDEp, NIEp), digits = 3))
208
209 rownames(parMedResults) <- c('ATE', 'NDE', 'NIE')
210
211 print(parMedResults)
```

# Example: NLSY79

- Compute triply robust estimates of the *NDE* and *NIE* using parametric models

```
> print(parMedResults)
      point.est      se ll.95ci ul.95ci pval
ATE      -0.150 0.022  -0.194  -0.107 0.00
NDE      -0.152 0.042  -0.234  -0.069 0.00
NIE       0.001 0.024  -0.045   0.048 0.95
> |
```

# Example: NLSY79

- Compute DML estimates of the *NDE* and *NIE* using super learners

```
213 # semi-parametric multiply robust estimate for NDE and NIE w/ cross-fitting #
214 set.seed(3308004)
215
216 nlsy$k <- runif(nrow(nlsy), 0, 1)
217 nlsy <- nlsy[order(nlsy$k), ]
218 nlsy$k <- rep(1:5, length.out=nrow(nlsy))
219
220 cntrl.sl <- SuperLearner.CV.control(V=10)
221
222 confounders <- c("female", "black", "hispan", "paredu",
223 "parprof", "parinc_prank", "famsize", "afqt3")
224
225 for (j in 1:5) {
226
227   df.train <- nlsy[which(nlsy$k!=j), ]
228   df.est <- nlsy[which(nlsy$k==j), ]
229
230   Ymodel.sl <- SuperLearner(
231     Y=df.train[, "std_cesd_age40"],
232     X=df.train[, c(confounders, "att22", "ever_unemp_age3539")],
233     SL.library=c("SL.lm", "SL.ranger", "SL.xgboost"),
234     cvControl=cntrl.sl)
235
236   Dmodel_1.sl <- SuperLearner(
237     Y=df.train[, "att22"],
238     X=df.train[, confounders],
239     family = binomial(),
240     SL.library=c("SL.glm", "SL.ranger", "SL.xgboost"),
241     cvControl=cntrl.sl)
242
243   Dmodel_2.sl <- SuperLearner(
244     Y=df.train[, "att22"],
245     X=df.train[, c(confounders, "ever_unemp_age3539")],
246     family = binomial(),
247     SL.library=c("SL.glm", "SL.ranger", "SL.xgboost"),
248     cvControl=cntrl.sl)
```

# Example: NLSY79

- Compute DML estimates of the *NDE* and *NIE* using super learners

```
250 phat_d1_C <- predict(Dmodel_1.sl,  
251   newdata=df.est[,confounders], type="prob")$pred  
252 phat_d0_C <- 1-phat_d1_C  
253  
254 phat_d1_CM <- predict(Dmodel_2.sl,  
255   newdata=df.est[,c(confounders, "ever_unemp_age3539")], type="prob")$pred  
256 phat_d0_CM <- 1-phat_d1_CM  
257  
258 phat_d1 <- mean(df.est$att22)  
259 phat_d0 <- 1-phat_d1  
260  
261 nlsy$f1[nlsy$k==j] <- nlsy$att22[nlsy$k==j]*phat_d1 / phat_d1_C  
262 nlsy$f2[nlsy$k==j] <- (1-nlsy$att22[nlsy$k==j])*phat_d0 / phat_d0_C  
263 nlsy$f3[nlsy$k==j] <- nlsy$att22[nlsy$k==j]*phat_d0 / phat_d0_C  
264  
265 nlsy$s1[nlsy$k==j] <- phat_d0_CM / phat_d1_CM  
266  
267 nlsy$resid[nlsy$k==j] <- df.est$std_cesd_age40 -  
268   predict(Ymodel.sl, df.est[,c(confounders, "att22", "ever_unemp_age3539")])$pred  
269  
270 gdata <- df.est[,c(confounders, "att22", "ever_unemp_age3539")]  
271  
272 gdata$att22 <- 1  
273 nlsy$qhat1[nlsy$k==j] <- predict(Ymodel.sl, newdata=gdata)$pred  
274  
275 gdata$att22 <- 0  
276 nlsy$qhat0[nlsy$k==j] <- predict(Ymodel.sl, newdata=gdata)$pred
```

# Example: NLSY79

- Compute DML estimates of the *NDE* and *NIE* using super learners

```
278 qhat0_CD <- lm(qhat0 ~ att22 * (female + black + hispan +
279   paredu + parprof + parinc_prank + famsize + afqt3),
280   data=nlsy[which(nlsy$k==j),])
281
282 gdata$att22 <- 0
283 nlsy$EMgivC0qhat0[nlsy$k==j] <- predict(qhat0_CD, newdata=gdata)
284
285 qhat1_CD <- lm(qhat1 ~ att22 * (female + black + hispan +
286   paredu + parprof + parinc_prank + famsize + afqt3),
287   data=nlsy[which(nlsy$k==j),])
288
289 gdata$att22 <- 1
290 nlsy$EMgivC1qhat1[nlsy$k==j] <- predict(qhat1_CD, newdata=gdata)
291
292 gdata$att22 <- 0
293 nlsy$EMgivC0qhat1[nlsy$k==j] <- predict(qhat1_CD, newdata=gdata)
294 }
295
296 eif_psi_11 <- nlsy$f1*1*nlsy$resid +
297   nlsy$f1*(nlsy$qhat1 - nlsy$EMgivC1qhat1) +
298   nlsy$EMgivC1qhat1
299
300 eif_psi_00 <- nlsy$f2*1*nlsy$resid +
301   nlsy$f2*(nlsy$qhat0 - nlsy$EMgivC0qhat0) +
302   nlsy$EMgivC0qhat0
303
304 eif_psi_10 <- nlsy$f3*nlsy$s1*nlsy$resid +
305   nlsy$f3*nlsy$s1*(nlsy$qhat1 - nlsy$EMgivC0qhat1) +
306   nlsy$EMgivC0qhat1
```

# Example: NLSY79

- Compute DML estimates of the *NDE* and *NIE* using super learners

```
308 NDEhat <- mean(eif_psi_10 - eif_psi_00)
309 NIEhat <- mean(eif_psi_11 - eif_psi_10)
310 ATEhat <- mean(eif_psi_11 - eif_psi_00)
311
312 VarNDEhat <- var(eif_psi_10 - eif_psi_00)/length(resid)
313 VarNIEhat <- var(eif_psi_11 - eif_psi_10)/length(resid)
314 VarATEhat <- var(eif_psi_11 - eif_psi_00)/length(resid)
315
316 NDE95pctCI <- c(NDEhat-1.96*sqrt(VarNDEhat), NDEhat+1.96*sqrt(VarNDEhat))
317 NIE95pctCI <- c(NIEhat-1.96*sqrt(VarNIEhat), NIEhat+1.96*sqrt(VarNIEhat))
318 ATE95pctCI <- c(ATEhat-1.96*sqrt(VarATEhat), ATEhat+1.96*sqrt(VarATEhat))
319
320 ATEp <- (1-pnorm(abs(ATEhat/sqrt(VarATEhat)),0,1))*2
321 NDEp <- (1-pnorm(abs(NDEhat/sqrt(VarNDEhat)),0,1))*2
322 NIEp <- (1-pnorm(abs(NIEhat/sqrt(VarNIEhat)),0,1))*2
323
324 sparMedResults <- data.frame(
325   point.est = round(c(ATEhat, NDEhat, NIEhat), digits = 3),
326   se = round(c(sqrt(VarATEhat), sqrt(VarNDEhat), sqrt(VarNIEhat)), digits = 3),
327   ll.95ci = round(c(ATE95pctCI[1], NDE95pctCI[1], NIE95pctCI[1]), digits = 3),
328   ul.95ci = round(c(ATE95pctCI[2], NDE95pctCI[2], NIE95pctCI[2]), digits = 3),
329   pval = round(c(ATEp, NDEp, NIEp), digits = 3))
330
331 rownames(sparMedResults) <- c('ATE', 'NDE', 'NIE')
332
333 print(sparMedResults)
```

# Example: NLSY79

- Compute DML estimates of the *NDE* and *NIE* using super learners

```
> print(sparMedResults)
      point.est      se ll.95ci ul.95ci  pval
ATE      -0.107 0.021  -0.147  -0.066 0.000
NDE      -0.115 0.038  -0.190  -0.039 0.003
NIE       0.008 0.022  -0.035   0.051 0.709
> |
```

# Limitations

- MR/DML estimation is best suited for applications with a binary exposure
  - With an exposure that is continuous or otherwise has many values, MR estimation becomes unstable and practically difficult to implement
- MR estimation can also perform poorly when *none* of the component models are correctly specified, and sometimes, the misspecification bias can be even worse than with fully parametric approaches
- When implemented with parametric models, it is safer to compute inferential statistics using the bootstrap



# Extensions

- Similar methods can also be used to construct MR/DML estimates for...
  - multivariate natural and path-specific effects (Miles et al. 2021, Zhou 2022)
  - interventional direct and indirect effects (Diaz et al. 2021, Rudolph et al. 2017)
- Targeted maximum likelihood (TMLE; Zheng and van der Laan 2012)
  - An alternative, two-step approach to MR estimation
- causal-Graphical Normalizing Flows (cGNFs; Balgi et al. 2024)
  - Learn entire causal systems represented as DAGs using highly flexible neural networks

# The problem of confounding


- Another important concern in analyses of causal mediation is the problem of unobserved confounding
- If there are unobserved confounders for the exposure-outcome, mediator-outcome, or exposure-mediator relationships, then all the estimators that we have discussed previously are biased and inconsistent
- Confounding bias poses a significant challenge in analyses of causal mediation because it is very difficult to conduct experiments that ensure the relationships among key variables are unconfounded by design

# Sensitivity analysis

- When confounding cannot be controlled by experimental design, formal sensitivity analyses are useful for assessing potential biases and their impact on our inferences
- A sensitivity analysis involves postulating different, hypothetical patterns of unobserved confounding and then exploring how the resulting bias may alter our conclusions about causal mediation

# Sensitivity analysis

- Consider a scenario in which an unobserved variable, denoted by  $U$ , affects the exposure  $D$ , the mediator  $M$ , and the outcome  $Y$
- Estimates of total, direct, and indirect effects would all be biased and inconsistent due to unobserved confounding



# Sensitivity analysis

- Suppose that...
  - the unobserved confounder  $U$  is binary
  - $E(Y|c, d, m, U = 1) - E(Y|c, d, m, U = 0)$  is constant in  $c$ ,  $d$ , and  $m$
  - $P(U = 1|c, d, m) - P(U = 1|c, d^*, m)$  is constant in  $c$  and  $m$
- In this scenario, the confounding bias in an estimator for the natural direct effect can be expressed as follows:

$$\begin{aligned} \text{Bias} \left( \widehat{NDE}(d, d^*) \right) &= [E(Y|c, d, m, U = 1) - E(Y|c, d, m, U = 0)] \times [P(U = 1|c, d, m) - P(U = 1|c, d^*, m)] \\ &= \delta\phi \end{aligned}$$

# Sensitivity analysis

- The confounding bias in an estimator for the natural direct effect:

$$\text{Bias}(\widehat{NDE}(d, d^*)) = \delta\phi$$

- $\delta = E(Y|c, d, m, U = 1) - E(Y|c, d, m, U = 0)$  is the difference in the mean of the outcome associated with a unit increase in the unobserved confounder, conditional on the baseline confounders and mediator
- $\phi = P(U = 1|c, d, m) - P(U = 1|c, d^*, m)$  is the difference in the probability of the unobserved confounder comparing level  $d$  versus  $d^*$  of the exposure, conditional on the confounders and mediator

# Sensitivity analysis

- Suppose further that...
  - $E(Y|c, d, U = 1) - E(Y|c, d, U = 0)$  is constant in  $c$  and  $d$
  - $P(U = 1|c, d) - P(U = 1|c, d^*)$  is constant in  $c$
- In this scenario, the confounding bias in an estimator for the natural indirect effect can be expressed as follows:

$$\begin{aligned} \text{Bias}\left(\widehat{NIE}(d, d^*)\right) &= [E(Y|c, d, U = 1) - E(Y|c, d, U = 0)] \times [P(U = 1|c, d) - P(U = 1|c, d^*)] - \delta\phi \\ &= \eta\psi - \delta\phi \end{aligned}$$

# Sensitivity analysis

- The confounding bias in an estimator for the natural indirect effect:

$$\text{Bias}(\widehat{NIE}(d, d^*)) = \eta\psi - \delta\phi$$

- $\delta$  and  $\phi$  are defined as before
- $\eta$  is the difference in the mean of the outcome associated with a unit increase in the unobserved confounder, conditional only on the baseline confounders and the exposure
- $\psi$  is the difference in the probability of the unobserved confounder comparing level  $d$  versus  $d^*$  of the exposure, conditional only on the baseline confounders



# Sensitivity analysis


- Under the same suppositions outlined previously...

$$\begin{aligned} \text{Bias} \left( \widehat{ATE}(d, d^*) \right) &= \text{Bias} \left( \widehat{NDE}(d, d^*) \right) + \text{Bias} \left( \widehat{NIE}(d, d^*) \right) \\ &= \delta\phi + (\eta\psi - \delta\phi) \\ &= \eta\psi \end{aligned}$$

where  $\eta = E(Y|c, d, U = 1) - E(Y|c, d, U = 0)$  and  $\psi = P(U = 1|c, d) - P(U = 1|c, d^*)$

# Sensitivity analysis

- Now consider a scenario in which an unobserved variable  $U$  affects only the mediator  $M$  and outcome  $Y$
- Estimates of direct and indirect effects would be biased and inconsistent due to unobserved confounding
- Estimates of total effects, however, would still be unbiased and consistent
- This scenario is common in standard experiments, where only the exposure is randomized



# Sensitivity Analysis

- In this scenario, and under the same assumptions about  $U$  as outlined previously, the bias in estimates of natural direct and indirect effects is given by the following expressions:

$$Bias\left(\widehat{NDE}(d, d^*)\right) = \delta\phi$$

$$Bias\left(\widehat{NIE}(d, d^*)\right) = -\delta\phi$$

where  $\delta$  and  $\phi$  are defined as before

# Sensitivity Analysis

- With the bias formulas outlined previously, a sensitivity analysis proceeds by reevaluating a set of effect estimates across different hypothetical patterns of unobserved confounding
- To this end, we evaluate the bias formulas across a range of plausible values for their sensitivity parameters, and then construct bias-adjusted effect estimates by subtracting the bias from the corresponding point estimate:

$$\widehat{NDE}(d, d^*)^{adj} = \widehat{NDE}(d, d^*) - Bias\left(\widehat{NDE}(d, d^*)\right)$$

$$\widehat{NIE}(d, d^*)^{adj} = \widehat{NIE}(d, d^*) - Bias\left(\widehat{NIE}(d, d^*)\right)$$

$$\widehat{ATE}(d, d^*)^{adj} = \widehat{ATE}(d, d^*) - Bias\left(\widehat{ATE}(d, d^*)\right)$$

# Example: NLSY79

- Compute bias-adjusted estimates for the natural direct and indirect effects of college attendance on depression, as mediated by unemployment

```
337 # sensitivity analysis for unobserved M-Y confounding #
338 sens.grid <- expand.grid(delta=seq(-0.2, 0.2, 0.02), phi=seq(-0.2, 0.2, 0.02))
339
340 adj.grid <- cbind(sens.grid,
341   nde.adj=sparMedResults[2,1]-(sens.grid$delta*sens.grid$phi),
342   nie.adj=sparMedResults[3,1]+(sens.grid$delta*sens.grid$phi))
343
```

# Example: NLSY79

- Compute bias-adjusted estimates for the NDE and NIE

```
344 nde.plot <- ggplot(adj.grid, aes(x=delta, y=phi, z=nde.adj, colour=stat(level))) +
345   geom_contour(
346     breaks=seq(round(min(adj.grid$nde.adj), 3),
347               round(max(adj.grid$nde.adj), 3), 0.005), show.legend=FALSE) +
348   scale_colour_distiller(palette="Greys", direction=1) +
349   xlab(expression(delta)) +
350   ylab(expression(phi)) +
351   scale_x_continuous(breaks=seq(-0.2, 0.2, 0.04)) +
352   scale_y_continuous(breaks=seq(-0.2, 0.2, 0.04)) +
353   ggtitle("A. Bias-adjusted NDE(1,0) Estimates") +
354   theme_bw(base_size=11) +
355   theme(
356     panel.grid.major=element_blank(),
357     panel.grid.minor=element_blank()) +
358   geom_text_contour(
359     breaks=seq(
360       round(min(adj.grid$nde.adj), 3),
361       round(max(adj.grid$nde.adj), 3), 0.005),
362     stroke=0.3,
363     size=3,
364     skip=0,
365     color="black")
```

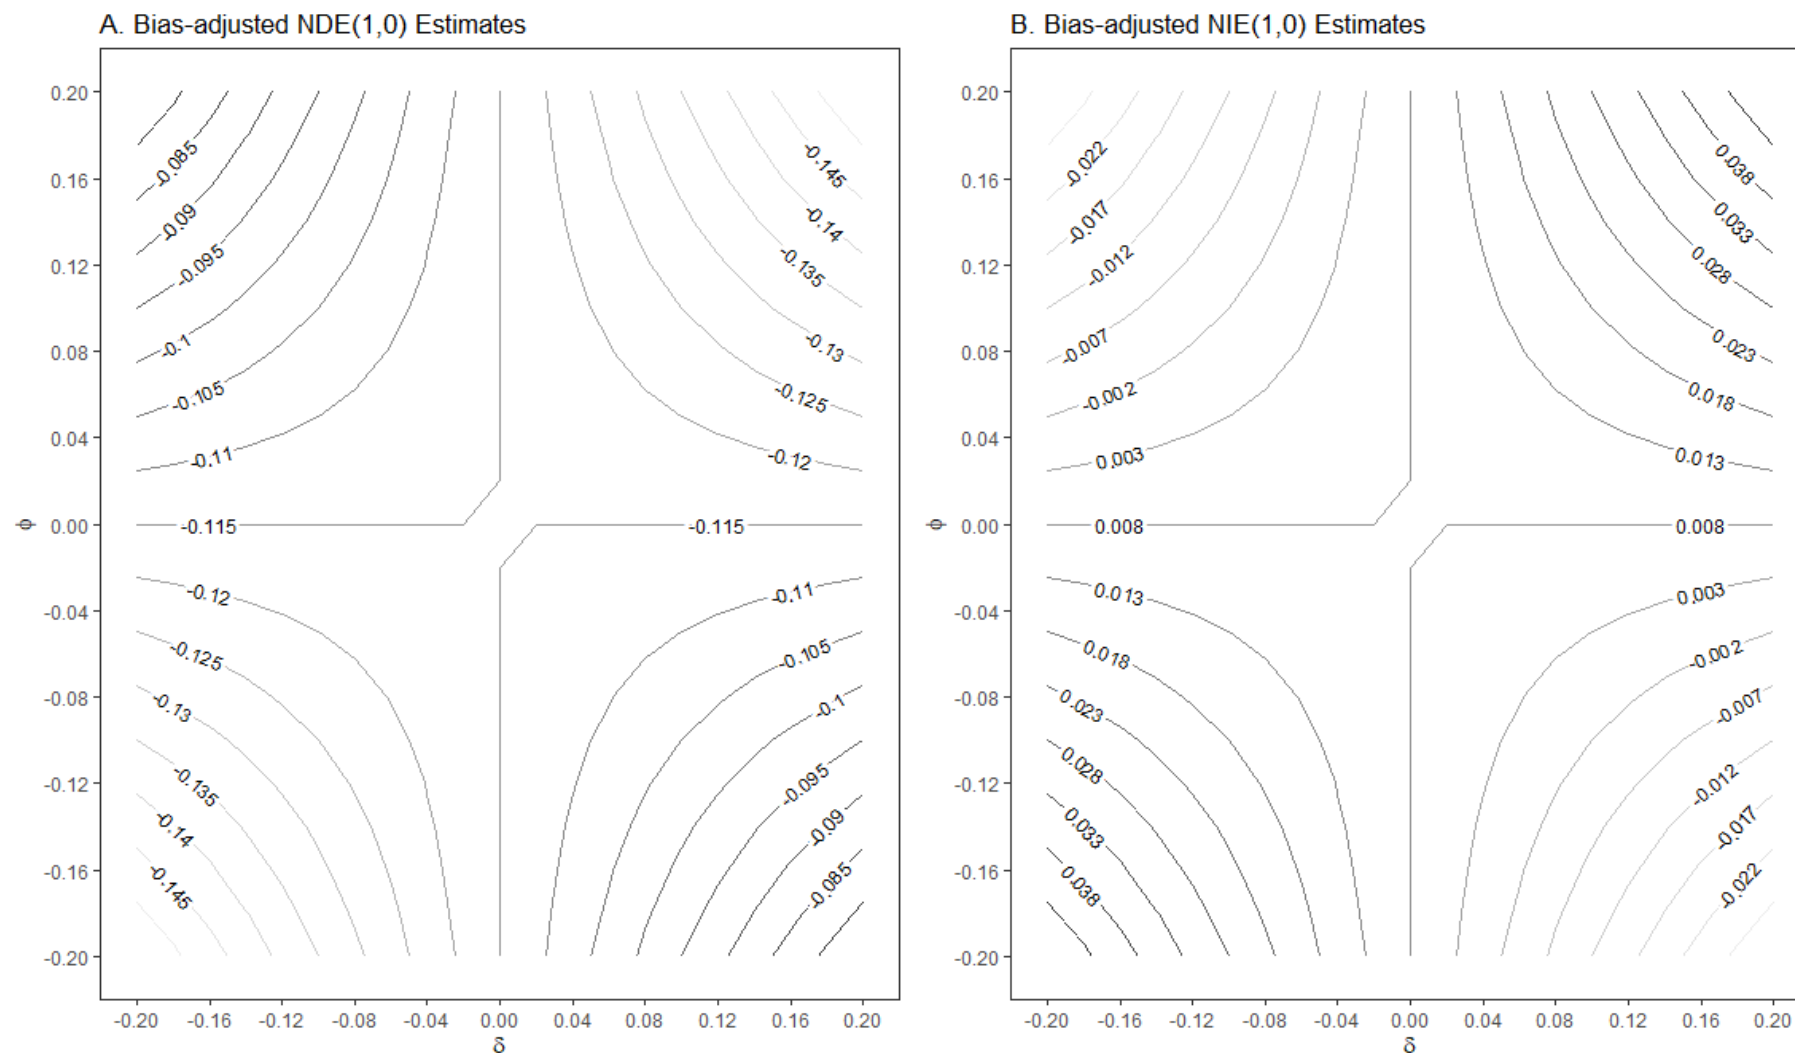
# Example: NLSY79

- Compute bias-adjusted estimates for the NDE and NIE

```
367 nie.plot <- ggplot(adj.grid, aes(x=delta,y=phi,z=nie.adj,colour=stat(level))) +
368   geom_contour(
369     breaks=seq(round(min(adj.grid$nue.adj),3),
370               round(max(adj.grid$nue.adj),3),0.005),show.legend=FALSE) +
371   scale_colour_distiller(palette="Greys",direction=1) +
372   xlab(expression(delta)) +
373   ylab(expression(phi)) +
374   scale_x_continuous(breaks=seq(-0.2,0.2,0.04)) +
375   scale_y_continuous(breaks=seq(-0.2,0.2,0.04)) +
376   ggtitle("B. Bias-adjusted NIE(1,0) Estimates") +
377   theme_bw(base_size=11) +
378   theme(
379     panel.grid.major=element_blank(),
380     panel.grid.minor=element_blank()) +
381   geom_text_contour(
382     breaks=seq(
383       round(min(adj.grid$nue.adj),3),
384       round(max(adj.grid$nue.adj),3),0.005),
385     stroke=0.3,
386     size=3,
387     skip=0,
388     color="black")
389
390 comb.plot <- grid.arrange(nde.plot, nie.plot, ncol=2, nrow=1)
391
392 print(comb.plot)
```

# Example: NLSY79

- Compute bias-adjusted estimates for the NDE and NIE





# Extensions

- Similar methods can also be used to construct bias-adjusted estimates for...
  - multivariate natural and path-specific effects (Zhou 2022)
  - interventional direct and indirect effects (Wodtke and Zhou 2020)